




A. CPU peak RSS vs T (full-batch) — $\sim O(T)$



B. CPU peak RSS vs chunk_size (T=785k) — bounded



C. GPU runtime + VRAM vs T (C=64) — $\sim O(T)$



D. GPU runtime + VRAM vs n_components (T=785k)

